Linda Petrini

AI safety research communication • writing, editing, research management
lindapetrini@gmail.com | +39 3516650672 | LinkedIn | Scholar

SKILLS

COMMUNICATION

Technical writing • Editing
Papers, blog posts, reports
Figures & data visualisation
Threads & newsletters
Video scripts • Talks

RESEARCH BACKGROUND

Interpretability • LLM evals
Alignment • Red teaming
Python • PyTorch • JAX • LaTeX

LANGUAGES

Italian (native) • English (fluent)
Russian (conversational)
French (conversational)

EDUCATION

MILA / UDEM

PhD, Computer Science (withdrew)
2021-2023 | Montreal, Canada
Deep learning interpretability for scientific discovery in epigenomics, concurrent with Google Brain.
Supervisors: M.G. Bellemare, A. Courville.

UNIVERSITY OF AMSTERDAM

MSc Artificial Intelligence
2017-2019 | Amsterdam, NL
Cum Laude.
Thesis: Generalization in representation learning via zero-shot learning.

UNIV. MILANO-BICOCCA

BSc Mathematics
2014-2017 | Milan, Italy

OUTREACH & SERVICE

Expert source on AI for El País
Public AI literacy workshop, Pontremoli (co-taught)
Reviewer: NeurIPS, ICML, ICLR
Mentor: AI4Good Lab (2019, 2020)
Co-organiser: CVPR workshop on Ethics in Creative Applications of Computer Vision

EXPERIENCE

ANTHROPIC | Research Manager, Anthropic Fellows Program

Sep 2026 - Present | Remote

- Research manager for a cohort of 20+ AI safety fellows, supporting research progress, career decisions and well-being through regular 1:1s.
- Writing and editing support for fellows' papers and posts ahead of publication.

AI SAFETY RESEARCH COMMUNICATION | Freelance

Sep 2023 - Present | Remote

- **Anthropic**: Writing, figures and editing for safety research papers and posts, including Alignment faking, SHADE-Arena, Constitutional classifiers and Inverse scaling in test-time compute.
- **Palisade Research**: Science communication for frontier AI risk evaluations. Drafted threads and figures (top thread: 1.3M views); shaped public framing of findings, and contributed to evaluation scoping, log review and red-teaming.
- **Foresight Institute**: Content lead for AI safety programmes. Built the Secure AI Tech Tree (concept, content, launch video); wrote AI Pathways scenarios and the Hyper Entities report; ran social media and newsletter (X: 11.4k); authored the AI for Climate and Nature workshop report with Bezos Earth Fund.
- **Epoch AI**: Editor for research reports, working directly with authors.
- **Nebius Science**: Researcher-facing articles on recent ML papers; hosted the weekly research seminar.
- **GSMA**: Co-author and editor on LLM benchmark papers for African languages.

GOOGLE BRAIN | AI Resident, then Research Associate

Oct 2019 - Jan 2023 | Montreal / Zurich

- Converted from the residency to a full-time research role.
- Led a cross-disciplinary team of 4 ML researchers and 4 epigenetics experts applying interpretability to epigenomics, bridging two fields with different vocabularies and standards of evidence.
- Research on set representation learning, Transformer losses, and Optimal Transport for genomic data.

SELECTED PUBLICATIONS

RESEARCH

- **Locality and Compositionality in Zero-Shot Learning**, ICLR 2020 (first author)
- **Mutual Information heatmaps as a tool for interpretability**, NeurIPS 2019 Workshop (first author)
- **Can We Trust Adversarial Filtering? When the Judge Curates the Benchmark**, Workshop paper, 2026
- **gradsim: Differentiable simulation for system identification and visuomotor control**, ICLR 2021
- **Ethics and Creativity in Computer Vision**, 2021

WRITING & FIGURES FOR SAFETY RESEARCH

- **Alignment faking in large language models**, 2024
- **SHADE-Arena: Evaluating sabotage and monitoring in LLM agents**, 2025
- **Inverse scaling in test-time compute**, 2025
- **Cost-effective constitutional classifiers via representation re-use**, 2025